


Enhancing Performance Tracing and Debugging in Remote Deployments

September 18, 2023

Google Cloud



Contents

- 01** Know the Speakers
- 02** Introduction
- 03** Trace Logs
- 04** In/Out of Memory
- 05** Analysis
- 06** Demo
- 07** Questions!

Speakers



Alankrit Kharbanda

GCP Data Transfer



AJ Ortega

GCP Data Transfer



Introduction

Intro to tracing

Identifying the problem

- **Tracing in On-premises remote deployed containers.**
 - **Tracing across multiple instances**
- **Machine not accessible to developer.**
- **Any changes including setting logstash or alternates is not possible because we deploy a standalone binary.**
- **Problem is three fold :**
 - **Generating Traces**
 - **Moving Traces to the cloud**
 - **Analyzing Traces**
- **Get everything together - Metrics , Traces and Logs**

Existing Solutions

Cloud Tracing

- Great for tracing across GCP services and more
- Great for latency tracking and traces.

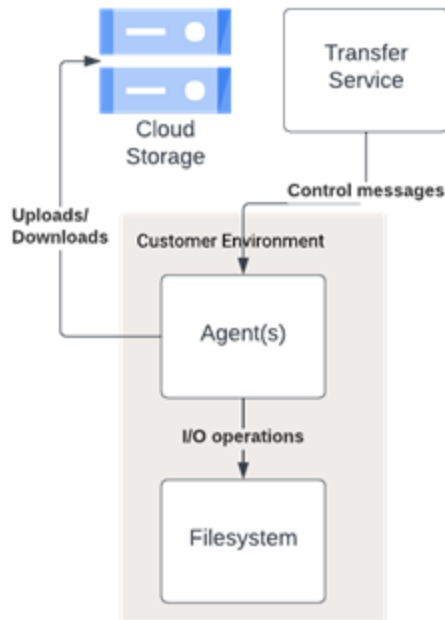
Other cloud tracing

- Powerful distributed tracing, headers in calls to other other cloud services.

- Both these services are great for Traces. Aggregation and analysis is easier with Bigquery using our method.

Solving Tracing in multi-container deployments

- Trace Beginning and end of methods
- Add metadata as key-value pairs
 - Metadata is tracked for current method/call.
 - Useful to add valuable side information.
 - Metadata can be any scalar type
- Attach ids to Context as “Spans”
 - Spans are tracked on subcalls as well
 - Useful in tracking sub-events to an event





Trace Logs

Sample code and Log entries

Code Sample(go)

```
1 func TestFirstMethod(t *testing.T) {
2     // mark start of the event
3     eventRecord := Start(ctx, "uber-event")
4
5     // defer the end of the event to the end of this method.
6     defer eventRecord.End(ctx)
7
8     // add metadata to the event.
9     eventRecord.set("test_key", "test_value")
10
11    // sub method call.
12    secondMethod(ctx)
13 }
14
15 func secondMethod(ctx context.Context) {
16     // mark start of the event
17     eventRecord := Start(ctx, "second_method")
18
19     // defer the end of the event to the end of this method.
20     defer eventRecord.End(ctx)
21
22     // add metadata to the event.
23     eventRecord.set("second_method_key", "second_method_value")
24
25     // do something in this method.
26     time.Sleep(2 * time.Second)
27 }
```

Event Sample (ends only)

```
1. {  
2.   "AgentID": "mock_agent_id",  
3.   "AgentPool": "mock_agent_pool",  
4.   "Duration": 2000,  
5.   "EndTime": 1692672280041,  
6.   "EventID": "2c14cb5c-c8cf-4712-94d6-8e64629f42a6",  
7.   "EventName": "second_method",  
8.   "EventPhase": "end",  
9.   "Metadata": {  
10.    "second_method_key": "second_method_value"  
11.  },  
12.   "ProjectID": "mock_project_name",  
13.   "Spans": {  
14.    "example-span-key": "example-span-value"  
15.  },  
16.   "Stack": [  
17.    "example-span-key"  
18.  ],  
19.   "StartTime": 1692672278041  
20. }  
21.
```

```
1. {  
2.   "AgentID": "mock_agent_id",  
3.   "AgentPool": "mock_agent_pool",  
4.   "Duration": 2001,  
5.   "EndTime": 1692672280041,  
6.   "EventID": "623425da-a159-4505-b7a7-95e95df110a5",  
7.   "EventName": "uber-event",  
8.   "EventPhase": "end",  
9.   "Metadata": {  
10.    "test_key": "test_value"  
11.  },  
12.   "ProjectID": "mock_project_name",  
13.   "Spans": {  
14.    "example-span-key": "example-span-value"  
15.  },  
16.   "Stack": [  
17.    "example-span-key"  
18.  ],  
19.   "StartTime": 1692672278040  
20. }  
21.
```



In stream vs out of stream

How to ship logs

Ship logs instream

- Google Cloud Logging
 - Use Cloud Logging SDK
 - Buffer Logs In Memory (Drop excess)
 - Pros :
 - No File Management
 - Easy SDK integration
 - Easy Integration with Bigquery for analysis
 - Cons :
 - Memory limitations by your application
 - Logs maybe dropped

Ship logs out of stream

- Google Cloud Storage

- Write Trace logs to file
- Separate threads to compress files
- Separate threads to send logs to GCS
- Pros :
 - Way less memory utilization, garbage generation.
 - Choose compression algorithm to optimize speed vs compressed file size.
 - Easy Integration with Bigquery for analysis
- Cons :
 - Delay in shipping and compressing files -> <5 minutes
 - Manually pull in for analysis.

Sampling - yes or no?

- Optional
- While sampling is great, our usecase was to solve for each trace
- Usecases where tracking all traces is helpful
 - Identifying events which never closed
 - Aggregation
 - Test statistics in aggregate
 - Hotspotting and concurrency trends
 - Pattern analysis for minority anomalies (eg. some small files causing problems vs most files are large).



Analysis

Lets import to Bigquery!

Analyze logs

- Bigquery

- Import logs

- Cloud Logging -> Direct Import Via Sink
- GCS Bucket -> Manual import -> Can import gz files.

- Queries

- Can save project relevant query templates directly
- Can write compounded queries easily -> Use temp tables if not great at sql to simplify problems
- Constant analysis, one time analysis.

- Visualize

- Create sample dashboards/charts. See Demo.




Real life problems we solved

Some problems we solved

- Identify long poles in subprocesses
 - Example - identified an issue in our os.Stat operation
 - OS level issue - would not have detected otherwise.
- Identify Concurrency and Hotspot issue
 - Identified that one of our systems was not honoring thread count limit
 - We were hammering a service we weren't supposed to.
- Close the loop measurement.
 - Identified that we were not releasing a lock that we were establishing.

Demo !

Google Cloud



Video

Thank you

